

Linear regression by using Deep Neural network: Implement Boston housing price prediction problem by Linear regression using Deep Neural network. Use Boston House price prediction dataset.

```
In [1]: 1 import io
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 %matplotlib inline
7 from sklearn.model_selection import train_test_split
8 from sklearn.preprocessing import StandardScaler
9 import tensorflow as tf
10 from tensorflow import keras
11 from tensorflow.keras import layers
12 from sklearn.metrics import mean_absolute_error, r2_score
13 import warnings
14 warnings.filterwarnings('ignore')
```

```
In [2]: 1 # Importing DataSet and take a Look at Data
2 data = pd.read_csv('housing_data - housing_data.csv')
3 data
```

```
Out[2]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	NaN	36.2
...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1	273	21.0	391.99	NaN	22.4
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1	273	21.0	396.90	9.08	20.6
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1	273	21.0	396.90	5.64	23.9
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1	273	21.0	393.45	6.48	22.0
505	0.04741	0.0	11.93	0.0	0.573	6.030	NaN	2.5050	1	273	21.0	396.90	7.88	11.9

506 rows × 14 columns

```
In [3]: 1 # Handle null values by filling them with the mean of the respective columns
2 data.fillna(data.mean(), inplace=True)
```

```
In [4]: 1 data.isnull().sum()
```

```
Out[4]: CRIM      0
ZN          0
INDUS       0
CHAS        0
NOX         0
RM          0
AGE         0
DIS         0
RAD         0
TAX         0
PTRATIO     0
B           0
LSTAT       0
MEDV        0
dtype: int64
```

```
In [5]: 1 data.describe()
```

Out[5]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.611874	11.211934	11.083992	0.069959	0.554695	6.284634	68.518519	3.795043	9.549407	408.237154	18.454382
std	8.545770	22.921051	6.699165	0.250233	0.115878	0.702617	27.439466	2.105710	8.707259	168.537116	2.160123
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.000000	12.600000
25%	0.083235	0.000000	5.190000	0.000000	0.449000	5.885500	45.925000	2.100175	4.000000	279.000000	17.400000
50%	0.290250	0.000000	9.900000	0.000000	0.538000	6.208500	74.450000	3.207450	5.000000	330.000000	19.050000
75%	3.611874	11.211934	18.100000	0.000000	0.624000	6.623500	93.575000	5.188425	24.000000	666.000000	20.200000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000	22.000000


```
In [6]: 1 data.info()
2 data.shape
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    CRIM         506 non-null    float64
1    ZN           506 non-null    float64
2    INDUS        506 non-null    float64
3    CHAS         506 non-null    float64
4    NOX          506 non-null    float64
5    RM           506 non-null    float64
6    AGE          506 non-null    float64
7    DIS          506 non-null    float64
8    RAD          506 non-null    int64
9    TAX          506 non-null    int64
10   PTRATIO      506 non-null    float64
11   B            506 non-null    float64
12   LSTAT        506 non-null    float64
13   MEDV         506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

Out[6]: (506, 14)


```
In [7]: 1 import seaborn as sns
2 sns.distplot(data.MEDV)
```

Out[7]: <AxesSubplot:xlabel='MEDV', ylabel='Density'>



```
In [8]: 1 sns.boxplot(data.MEDV)
```

```
Out[8]: <AxesSubplot:xlabel='MEDV'>
```



```
In [9]: 1 correlation = data.corr()  
2 correlation.loc['MEDV']
```

```
Out[9]: CRIM      -0.379695  
ZN          0.365943  
INDUS      -0.478657  
CHAS        0.179882  
NOX        -0.427321  
RM          0.695360  
AGE        -0.380223  
DIS         0.249929  
RAD        -0.381626  
TAX        -0.468536  
PTRATIO    -0.507787  
B           0.333461  
LSTAT      -0.721975  
MEDV       1.000000  
Name: MEDV, dtype: float64
```

```
In [10]: 1 # plotting the heatmap
2 import matplotlib.pyplot as plt
3 fig, axes = plt.subplots(figsize=(15,12))
4 sns.heatmap(correlation, square = True, annot = True)
```

Out[10]: <AxesSubplot:>



```
In [11]: 1 # Checking the scatter plot with the most correlated features
2 plt.figure(figsize = (20,5))
3 features = ['LSTAT', 'RM', 'PTRATIO']
4 for i, col in enumerate(features):
5     plt.subplot(1, len(features) , i+1)
6     x = data[col]
7     y = data.MEDV
8     plt.scatter(x, y, marker='o')
9     plt.title("Variation in House prices")
10    plt.xlabel(col)
11    plt.ylabel('House prices in $1000')
```



```
In [12]: 1 # Splitting the dependent feature and independent feature
2 #X = data[['LSTAT', 'RM', 'PTRATIO']]
3 X = data.iloc[:, :-1]
4 y= data.MEDV
```

```
In [13]: 1 import numpy as np
2 from sklearn.model_selection import train_test_split
3
4 # Assuming you have data stored in some variables X and y
5 # Splitting data into training and testing sets
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
7
8 # Now you can proceed with the code you provided
9 # Importing necessary libraries
10 from sklearn.linear_model import LinearRegression
11 from sklearn.preprocessing import StandardScaler
12
13 # Scaling the features
14 scaler = StandardScaler()
15 X_train_scaled = scaler.fit_transform(X_train)
16 X_test_scaled = scaler.transform(X_test)
17
18 mean = X_train.mean(axis=0)
19 std = X_train.std(axis=0)
20 X_train = (X_train - mean) / std
21 X_test = (X_test - mean) / std
22 #Linear Regression
23
24 from sklearn.linear_model import LinearRegression
25 regressor = LinearRegression()
26 #Fitting the model
27 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
28 regressor.fit(X_train,y_train)
```

```
Out[13]: ▾ LinearRegression
LinearRegression()
```

```
In [14]: 1 #Prediction on the test dataset
2 y_pred = regressor.predict(X_test)
3 # Predicting RMSE the Test set results
4 from sklearn.metrics import mean_squared_error
5 rmse = (np.sqrt(mean_squared_error(y_test, y_pred)))
6 print(rmse)
```

5.001766890194158

```
In [15]: 1 from sklearn.metrics import r2_score
2 r2 = r2_score(y_test, y_pred)
3 print(r2)
```

0.6588520195508162

In [16]:

```
1 import keras
2 from keras.layers import Dense
3 from keras.models import Sequential
4 from sklearn.preprocessing import StandardScaler
5 import matplotlib.pyplot as plt
6
7 # Assuming X_train and X_test are defined and initialized previously
8 # Assuming y_train is also defined and initialized
9
10 # Scaling the dataset
11 sc = StandardScaler()
12 X_train = sc.fit_transform(X_train)
13 X_test = sc.transform(X_test)
14
15 # Creating the neural network model
16 model = Sequential()
17 model.add(Dense(128, activation='relu', input_dim=13))
18 model.add(Dense(64, activation='relu'))
19 model.add(Dense(32, activation='relu'))
20 model.add(Dense(16, activation='relu'))
21 model.add(Dense(1))
22
23 # Compiling the model
24 model.compile(optimizer='adam', loss='mean_squared_error', metrics=['mae'])
25
26 # Visualizing the model architecture
27 keras.utils.plot_model(model, to_file='model.png', show_shapes=True, show_layer_names=True)
28
29 # Assuming you have defined your training data X_train and y_train
30 history = model.fit(X_train, y_train, epochs=100, validation_split=0.05)
31
32 # Plotting the training and validation Loss
33 plt.plot(history.history['loss'], label='Training Loss')
34 plt.plot(history.history['val_loss'], label='Validation Loss')
35 plt.xlabel('Epochs')
36 plt.ylabel('Loss')
37 plt.legend()
38 plt.show()
39
```

You must install pydot (`pip install pydot`) and install graphviz (see instructions at <https://graphviz.gitlab.io/download/>) (<https://graphviz.gitlab.io/download/>) for plot_model to work.

```
Epoch 1/100
12/12 [=====] - 2s 44ms/step - loss: 593.3365 - mae: 22.5031 - val_loss: 492.1255 - val_mae: 20.7200
Epoch 2/100
12/12 [=====] - 0s 9ms/step - loss: 525.0057 - mae: 20.9747 - val_loss: 399.8764 - val_mae: 18.5050
Epoch 3/100
12/12 [=====] - 0s 8ms/step - loss: 378.6923 - mae: 17.3772 - val_loss: 222.7657 - val_mae: 13.2538
Epoch 4/100
12/12 [=====] - 0s 10ms/step - loss: 155.3056 - mae: 10.3306 - val_loss: 52.9837 - val_mae: 4.9532
Epoch 5/100
12/12 [=====] - 0s 9ms/step - loss: 67.5092 - mae: 6.4644 - val_loss: 48.6647 - val_mae: 4.4646
Epoch 6/100
12/12 [=====] - 0s 9ms/step - loss: 37.0765 - mae: 4.5780 - val_loss: 56.6724 - val_mae: 4.6085
Epoch 7/100
12/12 [=====] - 0s 9ms/step - loss: 27.4871 - mae: 3.8358 - val_loss: 53.4959 - val_mae: 4.6053
Epoch 8/100
12/12 [=====] - 0s 9ms/step - loss: 22.7617 - mae: 3.5466 - val_loss: 52.5010 - val_mae: 4.6640
Epoch 9/100
12/12 [=====] - 0s 9ms/step - loss: 20.1987 - mae: 3.2901 - val_loss: 50.7301 - val_mae: 4.5293
Epoch 10/100
12/12 [=====] - 0s 10ms/step - loss: 18.5831 - mae: 3.1421 - val_loss: 49.0089 - val_mae: 4.3896
Epoch 11/100
12/12 [=====] - 0s 12ms/step - loss: 16.9307 - mae: 2.9703 - val_loss: 47.4259 - val_mae: 4.2535
Epoch 12/100
12/12 [=====] - 0s 9ms/step - loss: 15.8584 - mae: 2.8793 - val_loss: 44.7945 - val_mae: 4.1555
Epoch 13/100
12/12 [=====] - 0s 10ms/step - loss: 14.9110 - mae: 2.7769 - val_loss: 43.3636 - val_mae: 4.0691
Epoch 14/100
12/12 [=====] - 0s 10ms/step - loss: 14.2704 - mae: 2.7415 - val_loss: 42.3949 - val_mae: 4.0346
Epoch 15/100
12/12 [=====] - 0s 9ms/step - loss: 13.6585 - mae: 2.6691 - val_loss: 41.9805 - val_mae: 3.9858
Epoch 16/100
12/12 [=====] - 0s 7ms/step - loss: 13.3851 - mae: 2.6387 - val_loss: 39.0053 - val_mae: 3.8408
Epoch 17/100
12/12 [=====] - 0s 11ms/step - loss: 12.9651 - mae: 2.6520 - val_loss: 40.7240 - val_mae: 3.9047
Epoch 18/100
12/12 [=====] - 0s 11ms/step - loss: 12.5664 - mae: 2.5557 - val_loss: 40.6434 - val_mae: 3.9262
Epoch 19/100
12/12 [=====] - 0s 10ms/step - loss: 12.2979 - mae: 2.5747 - val_loss: 37.2956 - val_mae: 3.7524
Epoch 20/100
12/12 [=====] - 0s 7ms/step - loss: 11.8418 - mae: 2.4917 - val_loss: 37.6266 - val_mae: 3.7707
Epoch 21/100
12/12 [=====] - 0s 10ms/step - loss: 11.7158 - mae: 2.5176 - val_loss: 39.0760 - val_mae: 3.8018
Epoch 22/100
12/12 [=====] - 0s 9ms/step - loss: 11.4431 - mae: 2.4458 - val_loss: 38.5427 - val_mae: 3.6912
Epoch 23/100
12/12 [=====] - 0s 10ms/step - loss: 11.2924 - mae: 2.5149 - val_loss: 37.3225 - val_mae: 3.6865
Epoch 24/100
12/12 [=====] - 0s 9ms/step - loss: 11.1942 - mae: 2.4082 - val_loss: 36.1898 - val_mae: 3.6670
Epoch 25/100
12/12 [=====] - 0s 10ms/step - loss: 10.9826 - mae: 2.4827 - val_loss: 35.9758 - val_mae: 3.6670
```

l_mae: 3.6580
Epoch 26/100
12/12 [=====] - 0s 6ms/step - loss: 10.6356 - mae: 2.3610 - val_loss: 35.8442 - val_mae: 3.6747
Epoch 27/100
12/12 [=====] - 0s 8ms/step - loss: 10.5184 - mae: 2.4071 - val_loss: 35.2805 - val_mae: 3.5572
Epoch 28/100
12/12 [=====] - 0s 8ms/step - loss: 10.3016 - mae: 2.3593 - val_loss: 36.2390 - val_mae: 3.6871
Epoch 29/100
12/12 [=====] - 0s 9ms/step - loss: 10.3144 - mae: 2.3381 - val_loss: 33.6862 - val_mae: 3.6116
Epoch 30/100
12/12 [=====] - 0s 7ms/step - loss: 10.5037 - mae: 2.4386 - val_loss: 36.0907 - val_mae: 3.6182
Epoch 31/100
12/12 [=====] - 0s 10ms/step - loss: 9.8025 - mae: 2.3110 - val_loss: 33.4866 - val_mae: 3.6011
Epoch 32/100
12/12 [=====] - 0s 10ms/step - loss: 9.5271 - mae: 2.2879 - val_loss: 33.7885 - val_mae: 3.5193
Epoch 33/100
12/12 [=====] - 0s 9ms/step - loss: 9.3183 - mae: 2.2621 - val_loss: 33.1790 - val_mae: 3.5449
Epoch 34/100
12/12 [=====] - 0s 9ms/step - loss: 9.1131 - mae: 2.2614 - val_loss: 34.3521 - val_mae: 3.6106
Epoch 35/100
12/12 [=====] - 0s 9ms/step - loss: 8.8850 - mae: 2.2105 - val_loss: 32.9580 - val_mae: 3.5442
Epoch 36/100
12/12 [=====] - 0s 9ms/step - loss: 8.7797 - mae: 2.2195 - val_loss: 33.3213 - val_mae: 3.5653
Epoch 37/100
12/12 [=====] - 0s 10ms/step - loss: 8.6799 - mae: 2.2248 - val_loss: 31.7044 - val_mae: 3.4381
Epoch 38/100
12/12 [=====] - 0s 9ms/step - loss: 8.5635 - mae: 2.1858 - val_loss: 31.8526 - val_mae: 3.5358
Epoch 39/100
12/12 [=====] - 0s 10ms/step - loss: 8.4396 - mae: 2.1538 - val_loss: 32.1248 - val_mae: 3.5076
Epoch 40/100
12/12 [=====] - 0s 11ms/step - loss: 8.1019 - mae: 2.1395 - val_loss: 30.8459 - val_mae: 3.4163
Epoch 41/100
12/12 [=====] - 0s 11ms/step - loss: 7.9169 - mae: 2.0986 - val_loss: 31.0522 - val_mae: 3.4657
Epoch 42/100
12/12 [=====] - 0s 11ms/step - loss: 7.8935 - mae: 2.1001 - val_loss: 32.1896 - val_mae: 3.5412
Epoch 43/100
12/12 [=====] - 0s 11ms/step - loss: 7.7575 - mae: 2.1354 - val_loss: 30.1514 - val_mae: 3.3762
Epoch 44/100
12/12 [=====] - 0s 8ms/step - loss: 7.7905 - mae: 2.1055 - val_loss: 31.0291 - val_mae: 3.4593
Epoch 45/100
12/12 [=====] - 0s 9ms/step - loss: 7.3119 - mae: 2.0280 - val_loss: 29.6046 - val_mae: 3.4678
Epoch 46/100
12/12 [=====] - 0s 8ms/step - loss: 7.1641 - mae: 2.0162 - val_loss: 32.8597 - val_mae: 3.5630
Epoch 47/100
12/12 [=====] - 0s 8ms/step - loss: 7.0373 - mae: 1.9754 - val_loss: 29.0646 - val_mae: 3.3674
Epoch 48/100
12/12 [=====] - 0s 10ms/step - loss: 6.8431 - mae: 1.9862 - val_loss: 30.8615 - val_mae: 3.4752
Epoch 49/100
12/12 [=====] - 0s 8ms/step - loss: 6.8121 - mae: 1.9738 - val_loss: 28.5740 - val_mae: 3.3669
Epoch 50/100
12/12 [=====] - 0s 9ms/step - loss: 6.5312 - mae: 1.9086 - val_loss: 28.7225 - val_mae: 3.3653


Epoch 51/100
12/12 [=====] - 0s 9ms/step - loss: 6.4348 - mae: 1.9185 - val_loss: 28.8222 - val_mae: 3.3936
Epoch 52/100
12/12 [=====] - 0s 10ms/step - loss: 6.3505 - mae: 1.9192 - val_loss: 27.5490 - val_mae: 3.3383
Epoch 53/100
12/12 [=====] - 0s 9ms/step - loss: 6.0513 - mae: 1.8575 - val_loss: 27.0754 - val_mae: 3.3404
Epoch 54/100
12/12 [=====] - 0s 8ms/step - loss: 6.2330 - mae: 1.8712 - val_loss: 29.2607 - val_mae: 3.4202
Epoch 55/100
12/12 [=====] - 0s 9ms/step - loss: 5.9660 - mae: 1.8637 - val_loss: 28.4836 - val_mae: 3.4191
Epoch 56/100
12/12 [=====] - 0s 8ms/step - loss: 5.9944 - mae: 1.8457 - val_loss: 24.6683 - val_mae: 3.2299
Epoch 57/100
12/12 [=====] - 0s 9ms/step - loss: 5.6571 - mae: 1.8017 - val_loss: 27.9240 - val_mae: 3.4076
Epoch 58/100
12/12 [=====] - 0s 10ms/step - loss: 5.6048 - mae: 1.8299 - val_loss: 29.2053 - val_mae: 3.4728
Epoch 59/100
12/12 [=====] - 0s 9ms/step - loss: 5.4385 - mae: 1.7720 - val_loss: 23.6976 - val_mae: 3.1610
Epoch 60/100
12/12 [=====] - 0s 9ms/step - loss: 5.5935 - mae: 1.8099 - val_loss: 25.4217 - val_mae: 3.3188
Epoch 61/100

12/12 [=====] - 0s 9ms/step - loss: 5.3113 - mae: 1.7609 - val_loss: 27.5410 - val_mae: 3.4120
Epoch 62/100
12/12 [=====] - 0s 8ms/step - loss: 5.2426 - mae: 1.7450 - val_loss: 24.2072 - val_mae: 3.1952
Epoch 63/100
12/12 [=====] - 0s 8ms/step - loss: 5.2041 - mae: 1.7553 - val_loss: 25.6998 - val_mae: 3.3223
Epoch 64/100
12/12 [=====] - 0s 9ms/step - loss: 5.0989 - mae: 1.7075 - val_loss: 23.0396 - val_mae: 3.1827
Epoch 65/100
12/12 [=====] - 0s 9ms/step - loss: 5.1073 - mae: 1.7285 - val_loss: 26.8679 - val_mae: 3.3379
Epoch 66/100
12/12 [=====] - 0s 8ms/step - loss: 4.9866 - mae: 1.7151 - val_loss: 29.6811 - val_mae: 3.4808
Epoch 67/100
12/12 [=====] - 0s 8ms/step - loss: 5.0088 - mae: 1.7180 - val_loss: 22.0234 - val_mae: 3.0447
Epoch 68/100
12/12 [=====] - 0s 9ms/step - loss: 4.8893 - mae: 1.6920 - val_loss: 25.9152 - val_mae: 3.2718
Epoch 69/100
12/12 [=====] - 0s 9ms/step - loss: 4.5496 - mae: 1.6437 - val_loss: 25.2454 - val_mae: 3.2645
Epoch 70/100
12/12 [=====] - 0s 7ms/step - loss: 4.4986 - mae: 1.6102 - val_loss: 24.4274 - val_mae: 3.2512
Epoch 71/100
12/12 [=====] - 0s 7ms/step - loss: 4.3851 - mae: 1.5780 - val_loss: 23.7263 - val_mae: 3.1603
Epoch 72/100
12/12 [=====] - 0s 9ms/step - loss: 4.7020 - mae: 1.6812 - val_loss: 26.5144 - val_mae: 3.3664
Epoch 73/100
12/12 [=====] - 0s 10ms/step - loss: 4.4619 - mae: 1.6281 - val_loss: 25.6483 - val_mae: 3.2771
Epoch 74/100
12/12 [=====] - 0s 8ms/step - loss: 4.2446 - mae: 1.5621 - val_loss: 23.1437 - val_mae: 3.1621
Epoch 75/100
12/12 [=====] - 0s 7ms/step - loss: 4.1941 - mae: 1.5483 - val_loss: 24.4475 - val_mae: 3.2213
Epoch 76/100
12/12 [=====] - 0s 8ms/step - loss: 4.1007 - mae: 1.5233 - val_loss: 24.4011 - val_mae: 3.2016
Epoch 77/100
12/12 [=====] - 0s 7ms/step - loss: 3.9885 - mae: 1.5270 - val_loss: 22.9175 - val_mae: 3.1584
Epoch 78/100
12/12 [=====] - 0s 9ms/step - loss: 3.9726 - mae: 1.4958 - val_loss: 22.3622 - val_mae: 3.0758
Epoch 79/100
12/12 [=====] - 0s 8ms/step - loss: 3.9568 - mae: 1.5301 - val_loss: 24.2059 - val_mae: 3.2229
Epoch 80/100
12/12 [=====] - 0s 9ms/step - loss: 3.9104 - mae: 1.4846 - val_loss: 23.2432 - val_mae: 3.2169
Epoch 81/100
12/12 [=====] - 0s 7ms/step - loss: 3.7377 - mae: 1.4833 - val_loss: 24.3587 - val_mae: 3.2552
Epoch 82/100
12/12 [=====] - 0s 8ms/step - loss: 3.6685 - mae: 1.4626 - val_loss: 26.0521 - val_mae: 3.2781
Epoch 83/100
12/12 [=====] - 0s 7ms/step - loss: 3.7127 - mae: 1.4596 - val_loss: 23.1975 - val_mae: 3.1065
Epoch 84/100
12/12 [=====] - 0s 8ms/step - loss: 3.6200 - mae: 1.4455 - val_loss: 22.2413 - val_mae: 3.0606
Epoch 85/100
12/12 [=====] - 0s 7ms/step - loss: 3.5957 - mae: 1.4313 - val_loss: 22.5648 - val_mae: 3.1206
Epoch 86/100
12/12 [=====] - 0s 9ms/step - loss: 3.5110 - mae: 1.3960 - val_loss: 25.4949 - val_mae: 3.1206

```

mae: 3.2908
Epoch 87/100
12/12 [=====] - 0s 7ms/step - loss: 3.4958 - mae: 1.4208 - val_loss: 24.3938 - val_
mae: 3.1356
Epoch 88/100
12/12 [=====] - 0s 8ms/step - loss: 3.5582 - mae: 1.4469 - val_loss: 21.8915 - val_
mae: 3.0376
Epoch 89/100
12/12 [=====] - 0s 8ms/step - loss: 3.4020 - mae: 1.3666 - val_loss: 21.4599 - val_
mae: 2.9952
Epoch 90/100
12/12 [=====] - 0s 8ms/step - loss: 3.2326 - mae: 1.3574 - val_loss: 22.2129 - val_
mae: 3.1246
Epoch 91/100
12/12 [=====] - 0s 7ms/step - loss: 3.2493 - mae: 1.3628 - val_loss: 23.5151 - val_
mae: 3.1295
Epoch 92/100
12/12 [=====] - 0s 7ms/step - loss: 3.1388 - mae: 1.3269 - val_loss: 23.6589 - val_
mae: 3.0611
Epoch 93/100
12/12 [=====] - 0s 7ms/step - loss: 3.1233 - mae: 1.3342 - val_loss: 21.1233 - val_
mae: 2.9461
Epoch 94/100
12/12 [=====] - 0s 8ms/step - loss: 3.1661 - mae: 1.3309 - val_loss: 23.6154 - val_
mae: 3.1297
Epoch 95/100
12/12 [=====] - 0s 8ms/step - loss: 3.5468 - mae: 1.4289 - val_loss: 25.6488 - val_
mae: 3.1956
Epoch 96/100
12/12 [=====] - 0s 9ms/step - loss: 3.7170 - mae: 1.4744 - val_loss: 22.1357 - val_
mae: 3.1080
Epoch 97/100
12/12 [=====] - 0s 8ms/step - loss: 3.2998 - mae: 1.3654 - val_loss: 21.8552 - val_
mae: 3.0664
Epoch 98/100
12/12 [=====] - 0s 7ms/step - loss: 3.2627 - mae: 1.3816 - val_loss: 22.6286 - val_
mae: 3.0797
Epoch 99/100
12/12 [=====] - 0s 9ms/step - loss: 3.0908 - mae: 1.3260 - val_loss: 22.4770 - val_
mae: 3.0708
Epoch 100/100
12/12 [=====] - 0s 9ms/step - loss: 2.8196 - mae: 1.2830 - val_loss: 20.1691 - val_
mae: 2.8952

```



```

In [17]: 1 #Evaluation of the model
          2 y_pred = model.predict(X_test)
          3 mse_nn, mae_nn = model.evaluate(X_test, y_test)
          4 print('Mean squared error on test data: ', mse_nn)
          5 print('Mean absolute error on test data: ', mae_nn)

```

```

4/4 [=====] - 0s 4ms/step
4/4 [=====] - 0s 6ms/step - loss: 11.8925 - mae: 2.2676
Mean squared error on test data: 11.892463684082031
Mean absolute error on test data: 2.2675862312316895

```

```
In [18]: 1 #Comparison with traditional approaches
2 #First let's try with a simple algorithm, the Linear Regression:
3 from sklearn.metrics import mean_absolute_error
4 lr_model = LinearRegression()
5 lr_model.fit(X_train, y_train)
6 y_pred_lr = lr_model.predict(X_test)
7 mse_lr = mean_squared_error(y_test, y_pred_lr)
8 mae_lr = mean_absolute_error(y_test, y_pred_lr)
9 print('Mean squared error on test data: ', mse_lr)
10 print('Mean absolute error on test data: ', mae_lr)
11 from sklearn.metrics import r2_score
12 r2 = r2_score(y_test, y_pred)
13 print(r2)
```

Mean squared error on test data: 25.01767202384286
Mean absolute error on test data: 3.1499233573458034
0.8378310222849991

```
In [19]: 1 # Predicting RMSE the Test set results
2 from sklearn.metrics import mean_squared_error
3 rmse = (np.sqrt(mean_squared_error(y_test, y_pred)))
4 print(rmse)
```

3.448545294335442

```
In [20]: 1 # Make predictions on new data
2 import sklearn
3 new_data = scaler.transform([[0.1, 10.0, 5.0, 0, 0.4, 6.0, 50, 6.0, 1, 400, 20, 300, 10]]) # Scaling new
4 prediction = model.predict(new_data)
5 print("Predicted house price:", prediction)
```

1/1 [=====] - 0s 246ms/step
Predicted house price: [[14.101335]]